

0002182-lab4-gender-classification

April 5, 2023

Tải data vào colab

```
[85]: from google.colab import drive
drive.mount('/content/gdrive')
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

```
[86]: !unzip /content/gdrive/Shareddrives/nhan_dien_dac_diem_khuon_mat/gender_data.
      ↪ zip -d "/content"
```

Archive:

/content/gdrive/Shareddrives/nhan_dien_dac_diem_khuon_mat/gender_data.zip

replace /content/Training/female/131422.jpg.jpg? [y]es, [n]o, [A]ll, [N]one,

[r]ename: N

0.1 Import thư viện

```
[87]: import os
import cv2
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import glob
import seaborn as sn
from tqdm import tqdm
```

0.2 Chia 2 thư mục train và test

```
[88]: train_dir = '/content/Training'
test_dir = '/content/Validation'
categories = ["female", "male"]
img_size = 50
```

```
[89]: X_train_label = []
X_train_list = []
X_test_label = []
X_test_list = []
```

0.3 Load data và preprocessing

```
[90]: female_train_list = glob.glob(train_dir+"/"+categories[0]+"/*")
male_train_list = glob.glob(train_dir+"/"+categories[1]+"/*")
for name in female_train_list:
    X_train_label.append(0)
    img = cv2.imread(name)
    img = cv2.cvtColor(img,cv2.COLOR_BGR2GRAY)
    img = cv2.resize(img,(img_size,img_size))
    X_train_list.append((img))
for name in male_train_list:
    X_train_label.append(1)
    img = cv2.imread(name)
    img = cv2.cvtColor(img,cv2.COLOR_BGR2GRAY)
    img = cv2.resize(img,(img_size,img_size))
    X_train_list.append((img))
```

```
[ ]: print(len(X_train_list))
print(len(X_train_label))
```

47009

47009

```
[91]: female_test_list = glob.glob(test_dir+"/"+categories[0]+"/*")
male_test_list = glob.glob(test_dir+"/"+categories[1]+"/*")
for name in female_test_list:
    X_test_label.append(0)
    img = cv2.imread(name)
    img = cv2.cvtColor(img,cv2.COLOR_BGR2GRAY)
    img = cv2.resize(img,(img_size,img_size))
    X_test_list.append((img))
for name in male_test_list:
    X_test_label.append(1)
    img = cv2.imread(name)
    img = cv2.cvtColor(img,cv2.COLOR_BGR2GRAY)
    img = cv2.resize(img,(img_size,img_size))
    X_test_list.append((img))
```

```
[92]: print(len(X_test_list))
print(len(X_test_label))
```

11649

11649

```
[93]: X_train = np.array(X_train_list)
y_train = np.array(X_train_label)

print(X_train.shape)
```

```
print(y_train.shape)
```

```
(47009, 50, 50)
```

```
(47009,)
```

```
[94]: X_test = np.array(X_test_list)
      y_test = np.array(X_test_label)

      print(X_test.shape)
      print(y_test.shape)
```

```
(11649, 50, 50)
```

```
(11649,)
```

0.4 Duỗi vector, chuẩn hoá input

```
[95]: X_train_scaled = np.array([x.ravel()/255. for x in X_train])
      X_test_scaled = np.array([x.ravel()/255. for x in X_test])

      print(X_train_scaled.shape)
      print(X_test_scaled.shape)
```

```
(47009, 2500)
```

```
(11649, 2500)
```

0.5 Chuẩn hoá output

```
[96]: y_train_scaled = np.array([y for y in y_train])
      y_test_scaled = np.array([y for y in y_test])
```

0.6 Xây dựng và huấn luyện mô hình

```
[ ]: # Hàm sigmoid
def g(z):
    z = np.float64(z)
    return np.exp(z)/(1+np.exp(z))
# Hàm xác suất dự đoán
def predict_prob(X,w):
    z=np.dot(X,w)
    a = g(z)
    return a
# Hàm dự đoán binary
def predict(X,w):
    y_hat = predict_prob(X,w)
    y_hat[y_hat >=0.5] = 1
    y_hat[y_hat <0.5] = 0
    return y_hat
```

```

# Hàm loss
def loss(X,y,w):
    y_hat = predict_prob(X,w)
    theta=1e-4
    l = y*np.log(y_hat+theta) + (1-y)*np.log(1-y_hat+theta)
    return -np.mean(l)
# Hàm grad
def grad(X,y,w):
    y_hat = predict_prob(X,w)
    delta = y_hat - y
    dw = np.dot(X.T,delta)
    return dw
# Hàm gradient descent
def gradient_descent(X,y,lr = 1e-7, epochs = 1000):
    w = np.zeros((X.shape[1],1))
    losses = []
    for i in range(epochs):
        dw = grad(X,y,w)
        w -= lr*dw
        a = loss(X,y,w)
        losses.append(a)
    return losses,w


```

0.7 Vẽ đồ thị hàm loss ban đầu

```
[ ]: loss,w = gradient_descent(X_train_scaled,np.array([y_train_scaled]).T)
```

```
[ ]: plt.plot(loss)
plt.xlabel("epochs")
plt.ylabel("loss")
```

```
[ ]: Text(0, 0.5, 'loss')
```



```
[117]: print(f"Final loss: {loss[-1]}")
```

Final loss: 0.4909951632320063

```
[ ]: # Tính y dự đoán
y_hat = predict(X_test_scaled,w)
```

0.8 Hàm dự đoán

```
[ ]: def evaluation(y_true, y_pred):
    tp = tn = fp = fn = 0
    # So sánh y_true và y_pred
    for i in range(len(y_true)):
        if y_true[i] == 1 and y_pred[i] == 1:
            tp += 1
        elif y_true[i] == 0 and y_pred[i] == 0:
            tn += 1
        elif y_true[i] == 0 and y_pred[i] == 1:
            fp += 1
        elif y_true[i] == 1 and y_pred[i] == 0:
            fn += 1
```

```

accuracy = (tp+tn)/(tp+tn+fp+fn)
precision = (tp)/(tp+fp)
recall = (tp)/(tp+fn)
f1_score = 2*precision*recall/(precision+recall)
# Tính true positive, true negative, false positive, false negative
print(f"tp = {tp}, tn = {tn}, fp = {fp}, fn = {fn}")
# Tính Accuracy, Precision, Recall, F1 score
print(f"Accuracy = {accuracy}")
print(f"Precision = {precision}")
print(f"Recall = {recall}")
print(f"F1 Score = {f1_score}")

```

0.9 Đánh giá accuracy, recall, precision, f1-score

```
[ ]: evaluation(y_test, y_hat)
```

```

tp = 4491, tn = 4615, fp = 1226, fn = 1317
Accuracy = 0.7816979998283115
Precision = 0.785551862865139
Recall = 0.7732438016528925
F1 Score = 0.779349240780911

```

1 Ứng dụng phương pháp regularization L2 vào mô hình

```

[112]: # Hàm sigmoid
def g(z):
    z = np.float64(z)
    return np.exp(z)/(1+np.exp(z))
# Hàm xác suất dự đoán
def predict_prob(X, w):
    z = np.dot(X, w)
    a = g(z)
    return a
# Hàm dự đoán binary
def predict(X, w):
    y_hat = predict_prob(X, w)
    y_hat[y_hat >= 0.5] = 1
    y_hat[y_hat < 0.5] = 0
    return y_hat
# Hàm loss, có bổ sung L2 regularization
def loss(X, y, w, lambda_):
    y_hat = predict_prob(X, w)
    theta = 1e-4
    l1 = y*np.log(y_hat+theta) + (1-y)*np.log(1-y_hat+theta)
    l2 = lambda_ * np.sum(w**2)

```

```

    return -np.mean(l) + 12/(2*X.shape[0])
# Hàm gradient, có bổ sung L2 regularization
def grad(X, y, w, lambda_):
    y_hat = predict_prob(X, w)
    delta = y_hat - y
    dw = np.dot(X.T, delta)
    return dw
# Hàm gradient descent
def gradient_descent(X, y, lr=1e-7, epochs=1000, lambda_=1e-2):
    w = np.zeros((X.shape[1], 1))
    losses = []
    for i in range(epochs):
        dw = grad(X, y, w, lambda_)
        w -= lr * dw
        a = loss(X, y, w, lambda_)
        losses.append(a)
    return losses, w


```

1.1 Vẽ đồ thị Loss mới

```
[113]: loss, w = gradient_descent(X_train_scaled, np.array([y_train_scaled]).T)
```

```
[ ]: plt.plot(loss)
plt.xlabel("epochs")
plt.ylabel("loss")
```

```
[ ]: Text(0, 0.5, 'loss')
```



```
[123]: print(f"Final loss: {loss[-1]}")
```

Final loss: 0.4909951632320063

```
[124]: # Tính y dự đoán
y_hat = predict(X_test_scaled,w)
```

1.1.1 Tính Accuracy, recall, precision, f1-score

```
[129]: evaluation(y_test, y_hat)
```


```
tp = 4491, tn = 4615, fp = 1226, fn = 1317
Accuracy = 0.7816979998283115
Precision = 0.785551862865139
Recall = 0.7732438016528925
F1 Score = 0.779349240780911
```


1.2 Vẽ confusion matrix dùng thư viện sklearn

```
[ ]: import sklearn.metrics

# Tạo confusion matrix
cm = sklearn.metrics.confusion_matrix(y_test, y_hat)

# Vẽ confusion matrix
plt.figure()
sn.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```



1.3 Code và thực thi hàm batch_generator tối ưu hàm loss

```
[97]: def batch_generator(X, y, batch_size):
    # Chia nhiều batches để huấn luyện
    idx = np.arange(X.shape[0])
    np.random.shuffle(idx)

    n_batch = len(idx) // batch_size
    for i in range(n_batch):
        i_start = i * batch_size
        i_stop = min((i + 1) * batch_size, len(idx))
        yield X[idx[i_start:i_stop], :], y[idx[i_start:i_stop],:]

# Hàm sigmoid
def g(z):
    z = np.float64(z)
    return np.exp(z) / (1 + np.exp(z))

# Hàm xác suất dự đoán
def predict_prob(X, w):
    z = np.dot(X, w)
    a = g(z)
    return a

# Hàm dự đoán binary
def predict(X, w):
    y_hat = predict_prob(X, w)
    y_hat[y_hat >= 0.5] = 1
    y_hat[y_hat < 0.5] = 0
    return y_hat

# Hàm loss, có bổ sung L2 regularization
def loss(X, y, w, lambda_):
    y_hat = predict_prob(X, w)
    theta = 1e-5
    l = y*np.log(y_hat+theta) + (1-y)*np.log(1-y_hat+theta)
    l2 = lambda_ * np.sum(w**2)
    return -np.mean(l) + l2/(2*X.shape[0])

# Hàm gradient, có bổ sung L2 regularization
def grad(X, y, w, lambda_):
    y_hat = predict_prob(X, w)
    delta = y_hat - y
    dw = np.dot(X.T, delta)
    return dw

# Hàm gradient descent
```

```

def gradient_descent(X, y, lr=1e-6, epochs=1000, lambda_=1e-2):
    w = np.zeros((X.shape[1], 1))
    losses = []
    for i in range(epochs):
        for X_batch, y_batch in batch_generator(X, y, batch_size):
            dw = grad(X_batch, y_batch, w, lambda_)
            w -= lr * dw
        a = loss(X, y, w, lambda_)
        losses.append(a)
        # print(losses[-1])
    return losses, w

batch_size = 50
lr = 1e-7
epochs = 1000
lambda_ = 1e-5
losses, weight = gradient_descent(X_train_scaled, np.array([y_train_scaled]).T,
    ↪lr, epochs, lambda_)

print("Final loss:", losses[-1])

```

Final loss: 0.4911840690638267

```


[126]: plt.plot(losses)
plt.xlabel("epoch")
plt.ylabel("loss")

```

```

[126]: Text(0, 0.5, 'loss')

```



```
[127]: y_pred = predict(X_test_scaled,weight)
```

```
[128]: evaluation(y_test_scaled,y_pred)
```

```
tp = 4491, tn = 4615, fp = 1226, fn = 1317
Accuracy = 0.7816979998283115
Precision = 0.785551862865139
Recall = 0.7732438016528925
F1 Score = 0.779349240780911
```

2 Hàm binary1 bổ sung bias

```
[102]: class binary1:
    def __init__(self, lr = 1e-6, epochs = 1000):
        self.lr = lr
        self.epochs = epochs
    # Hàm thêm bias vào tập X
    def addBias(self,X):
        bias = np.ones((X.shape[0],1))
        X = np.hstack((bias,X));
```

```

        return X
    # Hàm sigmoid
    def g(z):
        z = np.float64(z)
        return np.exp(z)/(1+np.exp(z))
    # Hàm xác suất dự đoán
    def predict_prob(self,X, w):
        z = np.dot(X, w)
        a = g(z)
        return a
    # Hàm dự đoán binary
    def predict(self,X, w):
        y_hat = self.predict_prob(X, w)
        y_hat[y_hat >= 0.5] = 1
        y_hat[y_hat < 0.5] = 0
        return y_hat
    # Hàm loss
    def loss(self,X, y, w, lambda_):
        y_hat = self.predict_prob(X, w)
        theta = 1e-4
        l = y*np.log(y_hat+theta) + (1-y)*np.log(1-y_hat+theta)
        return -np.mean(l)
    # Hàm gradient
    def grad(self,X, y, w, lambda_):
        y_hat = self.predict_prob(X, w)
        delta = y_hat - y
        dw = np.dot(X.T, delta)
        return dw
    # Hàm gradient descent
    def gradient_descent(self,X, y, lr=1e-7, epochs=1000, lambda_=0.01):
        w = np.zeros((X.shape[1], 1))
        losses = []
        for i in range(epochs):
            dw = self.grad(X, y, w, lambda_)
            w -= lr * dw
            a = self.loss(X, y, w, lambda_)
            losses.append(a)
        return losses, w

```

```

[103]: model = binary1()
X_train_scaled_bias = model.addBias(X_train_scaled)
X_test_scaled_bias = model.addBias(X_test_scaled)
loss,w = model.gradient_descent(X_train_scaled_bias,np.array([y_train_scaled]).
↪T)

```


```

[104]: plt.xlabel("epochs")
plt.ylabel("loss")

```

```
plt.plot(loss)
```

```
[104]: [<matplotlib.lines.Line2D at 0x7eff668d6c70>]
```



```
[105]: print(f"Final loss: {loss[-1]}")
```

Final loss: 0.49053260160519135

3 Đánh giá mô hình sau khi thêm bias

```
[106]: y_pred = model.predict(X_test_scaled_bias,w)
```

```
[107]: evaluation(y_test_scaled,y_pred)
```

tp = 4495, tn = 4625, fp = 1216, fn = 1313
Accuracy = 0.7828998197270152
Precision = 0.7870775696025214
Recall = 0.7739325068870524
F1 Score = 0.7804496918135254