

Activity 1: Nhận diện chữ số viết tay (MNIST Dataset)

Khai báo thư viện sử dụng

In []:

```
from tensorflow.keras.datasets import mnist
import numpy as np
from matplotlib import pyplot as plt
import tensorflow
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Dropout, Flatten, Input, Reshape
from tensorflow.keras.layers import Conv2D, MaxPooling2D
```

Load data

In []:

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
print(X_train.shape, X_test.shape)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 [=====] - 0s 0us/step
(60000, 28, 28) (10000, 28, 28)

Chuẩn hóa dữ liệu input

In []:

```
X_train_scaled = np.array(X_train)/255.
X_test_scaled = np.array(X_test)/255.
```

Chuẩn hóa dữ liệu output

In []:

```
# OnehotVector output
from sklearn.preprocessing import OneHotEncoder
encoder = OneHotEncoder()
encoder.fit(y_train.reshape(-1,1))
y_train = encoder.transform(y_train.reshape(-1,1)).toarray()
y_test = encoder.transform(y_test.reshape(-1,1)).toarray()
```

Xây dựng mô hình CNN ban đầu

In []:

```
# CNN model
inp = Input(shape = (28,28,1)) # input shape
cnn = Conv2D(filters = 8, kernel_size = 3, activation='relu')(inp)
pooling = MaxPooling2D(pool_size=(2,2))(cnn)
drop = Dropout(0.2)(pooling)

cnn = Conv2D(filters = 16, kernel_size = 4, activation='relu')(drop)
pooling = MaxPooling2D(pool_size=(2,2))(cnn)
drop = Dropout(0.2)(pooling)

cnn = Conv2D(filters = 32, kernel_size = 4, activation='relu')(drop)
pooling = MaxPooling2D(pool_size=(2,2))(cnn)

f = Flatten()(pooling)
```

```
fc1 = Dense(units = 32, activation = 'relu')(f)
fc2 = Dense(units = 16, activation = 'relu')(fc1)
out = Dense(units = 10, activation = 'softmax')(fc2)

model = Model(inputs = inp, outputs = out)
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 28, 28, 1)]	0
conv2d (Conv2D)	(None, 26, 26, 8)	80
max_pooling2d (MaxPooling2D)	(None, 13, 13, 8)	0
dropout (Dropout)	(None, 13, 13, 8)	0
conv2d_1 (Conv2D)	(None, 10, 10, 16)	2064
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 16)	0
dropout_1 (Dropout)	(None, 5, 5, 16)	0
conv2d_2 (Conv2D)	(None, 2, 2, 32)	8224
max_pooling2d_2 (MaxPooling2D)	(None, 1, 1, 32)	0
flatten (Flatten)	(None, 32)	0
dense (Dense)	(None, 32)	1056
dense_1 (Dense)	(None, 16)	528
dense_2 (Dense)	(None, 10)	170
Total params: 12,122		
Trainable params: 12,122		
Non-trainable params: 0		

Huấn luyện mô hình ban đầu

In []:

```
optimizer1 = tensorflow.keras.optimizers.Adam(learning_rate = 0.001)
model.compile(optimizer = optimizer1, loss='categorical_crossentropy',metrics = ['accuracy'])

history = model.fit(X_train_scaled,y_train,batch_size=64,
                    epochs = 50, validation_data = (X_test_scaled, y_test))
```

Epoch 1/50
938/938 [=====] - 23s 7ms/step - loss: 0.5440 - accuracy: 0.8188
- val_loss: 0.1270 - val_accuracy: 0.9635
Epoch 2/50
938/938 [=====] - 6s 6ms/step - loss: 0.1699 - accuracy: 0.9472
- val_loss: 0.0851 - val_accuracy: 0.9745
Epoch 3/50
938/938 [=====] - 7s 7ms/step - loss: 0.1259 - accuracy: 0.9606
- val_loss: 0.0794 - val_accuracy: 0.9747
Epoch 4/50
938/938 [=====] - 7s 8ms/step - loss: 0.1018 - accuracy: 0.9682
- val_loss: 0.0536 - val_accuracy: 0.9838
Epoch 5/50
938/938 [=====] - 5s 5ms/step - loss: 0.0876 - accuracy: 0.9725
- val_loss: 0.0475 - val_accuracy: 0.9857

```
val_loss: 0.0417 - val_accuracy: 0.9861
Epoch 6/50
938/938 [=====] - 5s 6ms/step - loss: 0.0788 - accuracy: 0.9751
- val_loss: 0.0417 - val_accuracy: 0.9861
Epoch 7/50
938/938 [=====] - 5s 5ms/step - loss: 0.0704 - accuracy: 0.9779
- val_loss: 0.0379 - val_accuracy: 0.9881
Epoch 8/50
938/938 [=====] - 5s 6ms/step - loss: 0.0671 - accuracy: 0.9790
- val_loss: 0.0373 - val_accuracy: 0.9878
Epoch 9/50
938/938 [=====] - 5s 5ms/step - loss: 0.0595 - accuracy: 0.9807
- val_loss: 0.0400 - val_accuracy: 0.9862
Epoch 10/50
938/938 [=====] - 5s 5ms/step - loss: 0.0580 - accuracy: 0.9817
- val_loss: 0.0353 - val_accuracy: 0.9878
Epoch 11/50
938/938 [=====] - 5s 6ms/step - loss: 0.0558 - accuracy: 0.9825
- val_loss: 0.0364 - val_accuracy: 0.9867
Epoch 12/50
938/938 [=====] - 5s 5ms/step - loss: 0.0542 - accuracy: 0.9823
- val_loss: 0.0381 - val_accuracy: 0.9878
Epoch 13/50
938/938 [=====] - 8s 8ms/step - loss: 0.0489 - accuracy: 0.9843
- val_loss: 0.0307 - val_accuracy: 0.9899
Epoch 14/50
938/938 [=====] - 6s 7ms/step - loss: 0.0503 - accuracy: 0.9840
- val_loss: 0.0285 - val_accuracy: 0.9903
Epoch 15/50
938/938 [=====] - 6s 6ms/step - loss: 0.0462 - accuracy: 0.9852
- val_loss: 0.0322 - val_accuracy: 0.9886
Epoch 16/50
938/938 [=====] - 5s 5ms/step - loss: 0.0464 - accuracy: 0.9854
- val_loss: 0.0302 - val_accuracy: 0.9894
Epoch 17/50
938/938 [=====] - 5s 5ms/step - loss: 0.0434 - accuracy: 0.9856
- val_loss: 0.0292 - val_accuracy: 0.9890
Epoch 18/50
938/938 [=====] - 5s 6ms/step - loss: 0.0435 - accuracy: 0.9862
- val_loss: 0.0326 - val_accuracy: 0.9898
Epoch 19/50
938/938 [=====] - 5s 6ms/step - loss: 0.0406 - accuracy: 0.9866
- val_loss: 0.0304 - val_accuracy: 0.9906
Epoch 20/50
938/938 [=====] - 6s 6ms/step - loss: 0.0410 - accuracy: 0.9866
- val_loss: 0.0261 - val_accuracy: 0.9920
Epoch 21/50
938/938 [=====] - 5s 5ms/step - loss: 0.0391 - accuracy: 0.9875
- val_loss: 0.0311 - val_accuracy: 0.9900
Epoch 22/50
938/938 [=====] - 6s 6ms/step - loss: 0.0400 - accuracy: 0.9870
- val_loss: 0.0307 - val_accuracy: 0.9890
Epoch 23/50
938/938 [=====] - 5s 5ms/step - loss: 0.0389 - accuracy: 0.9874
- val_loss: 0.0300 - val_accuracy: 0.9887
Epoch 24/50
938/938 [=====] - 5s 5ms/step - loss: 0.0358 - accuracy: 0.9885
- val_loss: 0.0311 - val_accuracy: 0.9884
Epoch 25/50
938/938 [=====] - 6s 6ms/step - loss: 0.0367 - accuracy: 0.9886
- val_loss: 0.0279 - val_accuracy: 0.9903
Epoch 26/50
938/938 [=====] - 5s 6ms/step - loss: 0.0363 - accuracy: 0.9881
- val_loss: 0.0271 - val_accuracy: 0.9920
Epoch 27/50
938/938 [=====] - 5s 6ms/step - loss: 0.0373 - accuracy: 0.9879
- val_loss: 0.0283 - val_accuracy: 0.9907
Epoch 28/50
938/938 [=====] - 5s 5ms/step - loss: 0.0341 - accuracy: 0.9884
- val_loss: 0.0287 - val_accuracy: 0.9903
Epoch 29/50
938/938 [=====] - 5s 6ms/step - loss: 0.0354 - accuracy: 0.9887
- val_loss: 0.0298 - val_accuracy: 0.9903
```

```

val_loss: 0.0290 - val_accuracy: 0.9903
Epoch 30/50
938/938 [=====] - 5s 5ms/step - loss: 0.0337 - accuracy: 0.9893
- val_loss: 0.0289 - val_accuracy: 0.9904
Epoch 31/50
938/938 [=====] - 8s 8ms/step - loss: 0.0333 - accuracy: 0.9889
- val_loss: 0.0295 - val_accuracy: 0.9902
Epoch 32/50
938/938 [=====] - 6s 6ms/step - loss: 0.0347 - accuracy: 0.9895
- val_loss: 0.0286 - val_accuracy: 0.9899
Epoch 33/50
938/938 [=====] - 5s 5ms/step - loss: 0.0342 - accuracy: 0.9890
- val_loss: 0.0293 - val_accuracy: 0.9892
Epoch 34/50
938/938 [=====] - 6s 6ms/step - loss: 0.0325 - accuracy: 0.9893
- val_loss: 0.0281 - val_accuracy: 0.9908
Epoch 35/50
938/938 [=====] - 5s 5ms/step - loss: 0.0327 - accuracy: 0.9897
- val_loss: 0.0275 - val_accuracy: 0.9910
Epoch 36/50
938/938 [=====] - 6s 6ms/step - loss: 0.0301 - accuracy: 0.9898
- val_loss: 0.0263 - val_accuracy: 0.9905
Epoch 37/50
938/938 [=====] - 5s 5ms/step - loss: 0.0309 - accuracy: 0.9895
- val_loss: 0.0313 - val_accuracy: 0.9900
Epoch 38/50
938/938 [=====] - 5s 5ms/step - loss: 0.0318 - accuracy: 0.9894
- val_loss: 0.0274 - val_accuracy: 0.9909
Epoch 39/50
938/938 [=====] - 6s 7ms/step - loss: 0.0317 - accuracy: 0.9898
- val_loss: 0.0288 - val_accuracy: 0.9910
Epoch 40/50
938/938 [=====] - 7s 7ms/step - loss: 0.0314 - accuracy: 0.9896
- val_loss: 0.0294 - val_accuracy: 0.9898
Epoch 41/50
938/938 [=====] - 7s 7ms/step - loss: 0.0321 - accuracy: 0.9896
- val_loss: 0.0285 - val_accuracy: 0.9914
Epoch 42/50
938/938 [=====] - 6s 6ms/step - loss: 0.0286 - accuracy: 0.9904
- val_loss: 0.0308 - val_accuracy: 0.9907
Epoch 43/50
938/938 [=====] - 5s 6ms/step - loss: 0.0313 - accuracy: 0.9900
- val_loss: 0.0316 - val_accuracy: 0.9906
Epoch 44/50
938/938 [=====] - 8s 8ms/step - loss: 0.0303 - accuracy: 0.9898
- val_loss: 0.0259 - val_accuracy: 0.9921
Epoch 45/50
938/938 [=====] - 6s 6ms/step - loss: 0.0311 - accuracy: 0.9897
- val_loss: 0.0269 - val_accuracy: 0.9919
Epoch 46/50
938/938 [=====] - 7s 7ms/step - loss: 0.0291 - accuracy: 0.9905
- val_loss: 0.0261 - val_accuracy: 0.9910
Epoch 47/50
938/938 [=====] - 6s 6ms/step - loss: 0.0305 - accuracy: 0.9902
- val_loss: 0.0271 - val_accuracy: 0.9918
Epoch 48/50
938/938 [=====] - 5s 5ms/step - loss: 0.0295 - accuracy: 0.9909
- val_loss: 0.0257 - val_accuracy: 0.9917
Epoch 49/50
938/938 [=====] - 5s 6ms/step - loss: 0.0293 - accuracy: 0.9905
- val_loss: 0.0268 - val_accuracy: 0.9919
Epoch 50/50
938/938 [=====] - 5s 5ms/step - loss: 0.0291 - accuracy: 0.9905
- val_loss: 0.0279 - val_accuracy: 0.9912

```

Trực quan hóa kết quả Accuracy và Loss trên tập Train và Test



In []:

```

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])

```

```
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc = 'upper left')
plt.show()
# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'text'], loc='upper left')
plt.show()
```



Lưu mô hình ban đầu, load mô hình đã lưu từ máy

In []:

```
model.save('modell.h5')
from tensorflow.keras.models import load_model
modell = load_model('/content/modell.h5')
```

Load ảnh và sử dụng mô hình đã huấn luyện để nhận diện

In []:

```
# import cv2
# img = cv2.imread('/content/so2.png')
# print(img.shape)
# plt.imshow(img)
```

In []:

```
# gray_img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
# img_new = cv2.resize(gray_img, (28,28))
# print(gray_img.shape)
# plt.imshow(img_new, cmap = 'gray')
```

In []:

```
# img_scaled = np.array([img_new/255.])
# print(img_scaled.shape)
# y_hat = modell.predict(img_scaled)
# print(y_hat)
# print(np.argmax(y_hat))
```

Tính y dự đoán từ mô hình ban đầu đã lưu

In []:

```
y_hat = modell.predict(X_test_scaled)
```

313/313 [=====] - 1s 2ms/step

In []:

```
print(y_hat.shape)
print(y_test.shape)
```

```
(10000, 10)
(10000, 10)
```

Lấy argmax của y dự đoán và y test

In []:

```
y_test = np.argmax(y_test, axis=1)
y_hat = np.argmax(y_hat, axis=1)
```

Activity 2

Sử dụng classification_report trong thư viện Sklearn đánh giá kết quả mô hình ban đầu dựa trên kết quả dự đoán tập test

In []:

```
from sklearn.metrics import classification_report
target_names = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']
print(classification_report(y_test, y_hat, target_names=target_names))
```

precision recall f1-score support

0	0.99	1.00	1.00	980
1	0.99	1.00	1.00	1135
2	1.00	0.98	0.99	1032
3	0.99	1.00	0.99	1010
4	1.00	0.99	0.99	982
5	0.98	0.99	0.99	892
6	1.00	0.99	0.99	958
7	0.99	0.99	0.99	1028
8	1.00	0.99	0.99	974
9	0.98	0.99	0.99	1009

accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000


Sử dụng Confusion_matrix trong thư viện Sklearn biểu diễn kết quả dự đoán trên tập test

In []:

```
import sklearn.metrics
import seaborn as sn

# Tạo confusion matrix
cm = sklearn.metrics.confusion_matrix(y_test, y_hat)

# Vẽ confusion matrix
plt.figure()
sn.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```



Load data, chuẩn hoá input

In []:

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
X_train_scaled = np.array(X_train)/255.  
X_test_scaled = np.array(X_test)/255.
```

Chuẩn hoá output

In []:

```
# OnehotVector output  
from sklearn.preprocessing import OneHotEncoder  
encoder = OneHotEncoder()  
encoder.fit(y_train.reshape(-1,1))  
y_train = encoder.transform(y_train.reshape(-1,1)).toarray()  
y_test = encoder.transform(y_test.reshape(-1,1)).toarray()
```

Thay đổi mạng CNN theo ý kiến của riêng mình sao cho kết quả tốt nhất

In []:

```
# Khai báo thư viện BatchNormalization  
from tensorflow.keras.layers import BatchNormalization  
  
inp = Input(shape = (28, 28, 1))  
  
cnn = Conv2D(filters = 32, kernel_size = 3, activation = 'relu', padding = 'same')(inp)  
# Thêm padding  
cnn = BatchNormalization()(cnn)  
cnn = Conv2D(filters = 32, kernel_size = 3, activation = 'relu', padding = 'same')(cnn)  
cnn = BatchNormalization()(cnn)  
pooling = MaxPooling2D(pool_size = (2, 2))(cnn)  
  
cnn = Conv2D(filters = 64, kernel_size = 3, activation = 'relu', padding = 'same')(pooling)  
cnn = BatchNormalization()(cnn)  
cnn = Conv2D(filters = 64, kernel_size = 3, activation = 'relu', padding = 'same')(cnn)  
cnn = BatchNormalization()(cnn)  
pooling = MaxPooling2D(pool_size = (2, 2))(cnn)  
  
cnn = Conv2D(filters = 128, kernel_size = 3, activation = 'relu', padding = 'same')(pooling)  
cnn = BatchNormalization()(cnn)  
cnn = Conv2D(filters = 128, kernel_size = 3, activation = 'relu', padding = 'same')(cnn)  
cnn = BatchNormalization()(cnn)  
pooling = MaxPooling2D(pool_size = (2, 2))(cnn)  
  
cnn = Conv2D(filters = 256, kernel_size = 3, activation = 'relu', padding = 'same')(pooling)  
cnn = BatchNormalization()(cnn)  
cnn = Conv2D(filters = 256, kernel_size = 3, activation = 'relu', padding = 'same')(cnn)  
cnn = BatchNormalization()(cnn)  
pooling = MaxPooling2D(pool_size = (2, 2))(cnn)  
  
f = Flatten()(pooling)  
  
fc1 = Dense(units = 512, activation = 'relu')(f)  
drop = Dropout(0.2)(fc1)  
fc2 = Dense(units = 256, activation = 'relu')(drop)  
fc3 = Dense(units = 128, activation = 'relu')(fc2)  
fc4 = Dense(units = 64, activation = 'relu')(fc3)  
fc5 = Dense(units = 32, activation = 'relu')(fc4)  
out = Dense(units = 10, activation = 'softmax')(fc5)  
  
model = Model(inputs=inp, outputs=out)  
model.summary()
```

Model: "model_9"

Layer (type)	Output Shape	Param #
=====		
input_14 (InputLayer)	[(None, 28, 28, 1)]	0

conv2d_43 (Conv2D)	(None, 28, 28, 32)	320
batch_normalization (Batch Normalization)	(None, 28, 28, 32)	128
conv2d_44 (Conv2D)	(None, 28, 28, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 28, 28, 32)	128
max_pooling2d_40 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_45 (Conv2D)	(None, 14, 14, 64)	18496
batch_normalization_2 (Batch Normalization)	(None, 14, 14, 64)	256
conv2d_46 (Conv2D)	(None, 14, 14, 64)	36928
batch_normalization_3 (Batch Normalization)	(None, 14, 14, 64)	256
max_pooling2d_41 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_47 (Conv2D)	(None, 7, 7, 128)	73856
batch_normalization_4 (Batch Normalization)	(None, 7, 7, 128)	512
conv2d_48 (Conv2D)	(None, 7, 7, 128)	147584
batch_normalization_5 (Batch Normalization)	(None, 7, 7, 128)	512
max_pooling2d_42 (MaxPooling2D)	(None, 3, 3, 128)	0
conv2d_49 (Conv2D)	(None, 3, 3, 256)	295168
batch_normalization_6 (Batch Normalization)	(None, 3, 3, 256)	1024
conv2d_50 (Conv2D)	(None, 3, 3, 256)	590080
batch_normalization_7 (Batch Normalization)	(None, 3, 3, 256)	1024
max_pooling2d_43 (MaxPooling2D)	(None, 1, 1, 256)	0
flatten_9 (Flatten)	(None, 256)	0
dense_35 (Dense)	(None, 512)	131584
dropout_43 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 256)	131328
dense_37 (Dense)	(None, 128)	32896
dense_38 (Dense)	(None, 64)	8256
dense_39 (Dense)	(None, 32)	2080
dense_40 (Dense)	(None, 10)	330

=====

Total params: 1,481,994
Trainable params: 1,480,074
Non-trainable params: 1,920

Huấn luyện mô hình sau khi thay đổi cấu trúc mạng CNN

In []:

```
optimizer1 = tensorflow.keras.optimizers.Adam(learning_rate = 0.001)
model.compile(optimizer = optimizer1, loss='categorical_crossentropy', metrics = ['accuracy'])

history = model.fit(X_train_scaled, y_train, batch_size=64,
                    epochs = 50, validation_data = (X_test_scaled, y_test))
```

Epoch 1/50

938/938 [=====] - 20s 13ms/step - loss: 0.1407 - accuracy: 0.9610 - val_loss: 0.0402 - val_accuracy: 0.9885

Epoch 2/50

938/938 [=====] - 11s 11ms/step - loss: 0.0539 - accuracy: 0.9865 - val_loss: 0.0574 - val_accuracy: 0.9839

Epoch 3/50

938/938 [=====] - 11s 12ms/step - loss: 0.0449 - accuracy: 0.9888 - val_loss: 0.0534 - val_accuracy: 0.9858

Epoch 4/50

938/938 [=====] - 11s 12ms/step - loss: 0.0401 - accuracy: 0.9904 - val_loss: 0.0550 - val_accuracy: 0.9871

Epoch 5/50

938/938 [=====] - 12s 13ms/step - loss: 0.0330 - accuracy: 0.9922 - val_loss: 0.0262 - val_accuracy: 0.9934

Epoch 6/50

938/938 [=====] - 11s 11ms/step - loss: 0.0279 - accuracy: 0.9931 - val_loss: 0.0399 - val_accuracy: 0.9908

Epoch 7/50

938/938 [=====] - 11s 12ms/step - loss: 0.0279 - accuracy: 0.9934 - val_loss: 0.0257 - val_accuracy: 0.9938

Epoch 8/50

938/938 [=====] - 11s 12ms/step - loss: 0.0241 - accuracy: 0.9946 - val_loss: 0.0424 - val_accuracy: 0.9916

Epoch 9/50

938/938 [=====] - 10s 11ms/step - loss: 0.0232 - accuracy: 0.9952 - val_loss: 0.0479 - val_accuracy: 0.9909

Epoch 10/50

938/938 [=====] - 11s 12ms/step - loss: 0.0221 - accuracy: 0.9951 - val_loss: 0.0419 - val_accuracy: 0.9920

Epoch 11/50

938/938 [=====] - 11s 11ms/step - loss: 0.0185 - accuracy: 0.9957 - val_loss: 0.0340 - val_accuracy: 0.9931

Epoch 12/50

938/938 [=====] - 11s 11ms/step - loss: 0.0183 - accuracy: 0.9965 - val_loss: 0.0398 - val_accuracy: 0.9913

Epoch 13/50

938/938 [=====] - 11s 12ms/step - loss: 0.0161 - accuracy: 0.9964 - val_loss: 0.0286 - val_accuracy: 0.9939

Epoch 14/50

938/938 [=====] - 11s 11ms/step - loss: 0.0129 - accuracy: 0.9969 - val_loss: 0.0374 - val_accuracy: 0.9924

Epoch 15/50

938/938 [=====] - 11s 12ms/step - loss: 0.0120 - accuracy: 0.9975 - val_loss: 0.0381 - val_accuracy: 0.9919

Epoch 16/50

938/938 [=====] - 11s 12ms/step - loss: 0.0179 - accuracy: 0.9962 - val_loss: 0.0358 - val_accuracy: 0.9929

Epoch 17/50

938/938 [=====] - 10s 11ms/step - loss: 0.0099 - accuracy: 0.9980 - val_loss: 0.0313 - val_accuracy: 0.9938

Epoch 18/50

938/938 [=====] - 11s 12ms/step - loss: 0.0119 - accuracy: 0.9976 - val_loss: 0.0355 - val_accuracy: 0.9933

Epoch 19/50

938/938 [=====] - 11s 11ms/step - loss: 0.0138 - accuracy: 0.9971 - val_loss: 0.0370 - val_accuracy: 0.9928

Epoch 20/50

```
938/938 [=====] - 11s 11ms/step - loss: 0.0107 - accuracy: 0.997
7 - val_loss: 0.0371 - val_accuracy: 0.9931
Epoch 21/50
938/938 [=====] - 11s 12ms/step - loss: 0.0102 - accuracy: 0.997
9 - val_loss: 0.0288 - val_accuracy: 0.9950
Epoch 22/50
938/938 [=====] - 11s 11ms/step - loss: 0.0097 - accuracy: 0.997
9 - val_loss: 0.0288 - val_accuracy: 0.9939
Epoch 23/50
938/938 [=====] - 11s 12ms/step - loss: 0.0085 - accuracy: 0.998
2 - val_loss: 0.0453 - val_accuracy: 0.9936
Epoch 24/50
938/938 [=====] - 11s 12ms/step - loss: 0.0108 - accuracy: 0.998
2 - val_loss: 0.0308 - val_accuracy: 0.9945
Epoch 25/50
938/938 [=====] - 10s 11ms/step - loss: 0.0073 - accuracy: 0.998
5 - val_loss: 0.0535 - val_accuracy: 0.9925
Epoch 26/50
938/938 [=====] - 11s 11ms/step - loss: 0.0079 - accuracy: 0.998
5 - val_loss: 0.0491 - val_accuracy: 0.9937
Epoch 27/50
938/938 [=====] - 11s 12ms/step - loss: 0.0098 - accuracy: 0.998
2 - val_loss: 0.0743 - val_accuracy: 0.9925
Epoch 28/50
938/938 [=====] - 11s 11ms/step - loss: 0.0102 - accuracy: 0.998
3 - val_loss: 0.0488 - val_accuracy: 0.9942
Epoch 29/50
938/938 [=====] - 11s 11ms/step - loss: 0.0065 - accuracy: 0.998
8 - val_loss: 0.0381 - val_accuracy: 0.9945
Epoch 30/50
938/938 [=====] - 11s 12ms/step - loss: 0.0058 - accuracy: 0.998
7 - val_loss: 0.0482 - val_accuracy: 0.9936
Epoch 31/50
938/938 [=====] - 11s 11ms/step - loss: 0.0066 - accuracy: 0.998
7 - val_loss: 0.0471 - val_accuracy: 0.9932
Epoch 32/50
938/938 [=====] - 10s 11ms/step - loss: 0.0078 - accuracy: 0.998
7 - val_loss: 0.0568 - val_accuracy: 0.9941
Epoch 33/50
938/938 [=====] - 11s 11ms/step - loss: 0.0069 - accuracy: 0.998
5 - val_loss: 0.0433 - val_accuracy: 0.9947
Epoch 34/50
938/938 [=====] - 11s 11ms/step - loss: 0.0072 - accuracy: 0.998
6 - val_loss: 0.0379 - val_accuracy: 0.9946
Epoch 35/50
938/938 [=====] - 11s 11ms/step - loss: 0.0046 - accuracy: 0.999
0 - val_loss: 0.0572 - val_accuracy: 0.9942
Epoch 36/50
938/938 [=====] - 11s 12ms/step - loss: 0.0068 - accuracy: 0.998
6 - val_loss: 0.0340 - val_accuracy: 0.9945
Epoch 37/50
938/938 [=====] - 12s 13ms/step - loss: 0.0061 - accuracy: 0.998
8 - val_loss: 0.0395 - val_accuracy: 0.9941
Epoch 38/50
938/938 [=====] - 11s 12ms/step - loss: 0.0043 - accuracy: 0.999
1 - val_loss: 0.0482 - val_accuracy: 0.9940
Epoch 39/50
938/938 [=====] - 11s 12ms/step - loss: 0.0078 - accuracy: 0.998
5 - val_loss: 0.0358 - val_accuracy: 0.9955
Epoch 40/50
938/938 [=====] - 11s 12ms/step - loss: 0.0041 - accuracy: 0.999
3 - val_loss: 0.0390 - val_accuracy: 0.9956
Epoch 41/50
938/938 [=====] - 11s 12ms/step - loss: 0.0060 - accuracy: 0.999
1 - val_loss: 0.0422 - val_accuracy: 0.9932
Epoch 42/50
938/938 [=====] - 11s 11ms/step - loss: 0.0063 - accuracy: 0.998
8 - val_loss: 0.0509 - val_accuracy: 0.9931
Epoch 43/50
938/938 [=====] - 11s 11ms/step - loss: 0.0069 - accuracy: 0.998
6 - val_loss: 0.0519 - val_accuracy: 0.9942
Epoch 44/50
```

```

938/938 [=====] - 11s 11ms/step - loss: 0.0057 - accuracy: 0.999
1 - val_loss: 0.0349 - val_accuracy: 0.9938
Epoch 45/50
938/938 [=====] - 11s 12ms/step - loss: 0.0049 - accuracy: 0.999
1 - val_loss: 0.0507 - val_accuracy: 0.9943
Epoch 46/50
938/938 [=====] - 11s 11ms/step - loss: 0.0043 - accuracy: 0.999
3 - val_loss: 0.0640 - val_accuracy: 0.9939
Epoch 47/50
938/938 [=====] - 12s 12ms/step - loss: 0.0044 - accuracy: 0.999
2 - val_loss: 0.0610 - val_accuracy: 0.9948
Epoch 48/50
938/938 [=====] - 11s 11ms/step - loss: 0.0048 - accuracy: 0.999
2 - val_loss: 0.1020 - val_accuracy: 0.9899
Epoch 49/50
938/938 [=====] - 11s 11ms/step - loss: 0.0080 - accuracy: 0.998
5 - val_loss: 0.0436 - val_accuracy: 0.9945
Epoch 50/50
938/938 [=====] - 11s 11ms/step - loss: 0.0024 - accuracy: 0.999
5 - val_loss: 0.0608 - val_accuracy: 0.9938

```

Trực quan hóa kết quả Accuracy và Loss trên tập Train và Test đối với mô hình mới

In []:


```

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc = 'upper left')
plt.show()
# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

```



model loss



Lưu mô hình mới

In []:

```
model.save('model2.h5')
from tensorflow.keras.models import load_model
model2 = load_model('/content/model2.h5')
```

Tính y dự đoán từ mô hình mới đã lưu, lấy argmax của y dự đoán và y test

In []:

```
y_hat = model2.predict(X_test_scaled)
y_test = np.argmax(y_test, axis=1)
y_hat = np.argmax(y_hat, axis=1)
```

313/313 [=====] - 2s 5ms/step

Sử dụng classification_report trong thư viện Sklearn đánh giá kết quả mô hình mới dựa trên kết quả dự đoán tập test

In []:

```
from sklearn.metrics import classification_report
target_names = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']
print(classification_report(y_test, y_hat, target_names=target_names))
```

	precision	recall	f1-score	support
0	0.99	1.00	1.00	980
1	0.99	1.00	0.99	1135
2	1.00	1.00	1.00	1032
3	0.99	1.00	0.99	1010
4	0.99	1.00	1.00	982
5	0.99	0.99	0.99	892
6	1.00	0.99	0.99	958
7	1.00	0.98	0.99	1028
8	0.99	0.99	0.99	974
9	1.00	0.99	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

Sử dụng Confusion_matrix trong thư viện Sklearn biểu diễn kết quả dự đoán trên tập test

In []:

```
import sklearn.metrics
import seaborn as sn

# Tạo confusion matrix
cm = sklearn.metrics.confusion_matrix(y_test, y_hat)

# Vẽ confusion matrix
plt.figure()
sn.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```

