

# uong-20002182-lab8-siamese-network

May 13, 2023

## 1 Siamese Network

Load data:

```
[ ]: from tensorflow.keras.datasets.mnist import load_data

(X_train, y_train), (X_test, y_test) = load_data()

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

```
11490434/11490434 [=====] - 0s 0us/step
(60000, 28, 28)
(10000, 28, 28)
(60000,)
```

Build model CNN

```
[ ]: from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Input
from tensorflow.keras.models import Model
from tensorflow.keras import backend as K

inp = Input(shape=(28,28,1))
x = Conv2D(filters = 8, kernel_size = 3, activation = 'relu')(inp)
x = MaxPooling2D(pool_size = (2,2))(x)

x = Flatten()(x)
x = Dense(units = 32, activation = 'relu')(x)
x = Dense(units = 2)(x)

cnn = Model(inputs = inp, outputs = x)
img1 = Input(shape = (28,28,1))
img2 = Input(shape = (28,28,1))

f1 = cnn(img1)
f2 = cnn(img2)
```

```

d = K.sqrt(K.sum(K.square(f1 - f2),axis = 1, keepdims = True))

model = Model(inputs = [img1,img2], outputs = d)
model.summary()
cnn.summary()

def loss(y_true, y_pred):
    proba = K.exp(-K.square(y_pred))
    return -K.mean(y_true * K.log(proba) + (1-y_true) * K.log(1-proba))

def loss1(y_true, y_pred):
    return K.mean(y_true * K.square(y_pred) + (1-y_true) * K.square(K.maximum(1.0,
↪- y_pred, 0)))

model.compile(optimizer = 'adam', loss = loss1)

```

Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, 28, 28, 1)]	0	[]
input_3 (InputLayer)	[(None, 28, 28, 1)]	0	[]
model (Functional) ['input_2[0][0]', 'input_3[0][0]']	(None, 2)	43442	
tf.math.subtract (TFOpLambda)	(None, 2)	0	['model[0][0]', 'model[1][0]']
tf.math.square (TFOpLambda) ['tf.math.subtract[0][0]']	(None, 2)	0	
tf.math.reduce_sum (TFOpLambda ['tf.math.square[0][0]'] )	(None, 1)	0	
tf.math.maximum (TFOpLambda) ['tf.math.reduce_sum[0][0]']	(None, 1)	0	
tf.math.sqrt (TFOpLambda) ['tf.math.maximum[0][0]']	(None, 1)	0	

```

=====
Total params: 43,442
Trainable params: 43,442
Non-trainable params: 0
-----

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 28, 28, 1)]	0
conv2d (Conv2D)	(None, 26, 26, 8)	80
max_pooling2d (MaxPooling2D)	(None, 13, 13, 8)	0
flatten (Flatten)	(None, 1352)	0
dense (Dense)	(None, 32)	43296
dense_1 (Dense)	(None, 2)	66

```

=====
Total params: 43,442
Trainable params: 43,442
Non-trainable params: 0
-----

```

Make all pairs or other strategies; some innovation here

```

[ ]: import numpy as np
from matplotlib import pyplot as plt

def generator(X,y,k = 8):
    unique_labels = np.unique(y)

    while True:
        X1 = []
        X2 = []
        y_batch = []
        for label in unique_labels:
            label_idx = np.where(y == label)[0]
            other_labels = set(unique_labels) - {label}

            for i in range(k):
                i1 = np.random.choice(label_idx)

```

```

        i2 = np.random.choice(label_idx)
        # i1 must be different from i2
        while i1 == i2:
            i2 = np.random.choice(label_idx)
        # create positive example
        X1.append(X[i1][:,:,None])
        X2.append(X[i2][:,:,None])
        y_batch.append(1.0)


        # create negative example
        i1 = np.random.choice(label_idx)
        my_label = np.random.choice(list(other_labels))
        i2 = np.random.choice(list(np.where(y == my_label)[0]))
        X1.append(X[i1][:,:,None])
        X2.append(X[i2][:,:,None])
        y_batch.append(0.0)
    yield [np.array(X1) / 255., np.array(X2) / 255.], np.array(y_batch)

# For testing
for pair, y in generator(X_test, y_test):
    print('Batch size: ', len(y))
    idx = np.random.choice(range(len(y)))
    print(pair[0][idx].shape)
    print('Pair label:', y[idx])
    plt.subplot(121)
    plt.imshow(pair[0][idx].reshape(28,28), cmap = 'gray')
    plt.subplot(122)
    plt.imshow(pair[1][idx].reshape(28,28), cmap = 'gray')

    break

```

Batch size: 160  
 (28, 28, 1)  
 Pair label: 1.0



### Fit model

```
[ ]: history = model.fit(generator(X_train, y_train, k = 32),
                           steps_per_epoch = 10,
                           epochs = 50,
                           validation_data = generator(X_test, y_test, k = 8),
                           validation_steps = 5)
```

```
Epoch 1/50
10/10 [=====] - 3s 279ms/step - loss: 0.0587 -
val_loss: 0.0643
Epoch 2/50
10/10 [=====] - 2s 274ms/step - loss: 0.0590 -
val_loss: 0.0679
Epoch 3/50
10/10 [=====] - 3s 293ms/step - loss: 0.0557 -
val_loss: 0.0654
Epoch 4/50
10/10 [=====] - 3s 361ms/step - loss: 0.0554 -
val_loss: 0.0645
Epoch 5/50
10/10 [=====] - 2s 268ms/step - loss: 0.0572 -
val_loss: 0.0649
Epoch 6/50
10/10 [=====] - 2s 272ms/step - loss: 0.0590 -
val_loss: 0.0615
Epoch 7/50
10/10 [=====] - 2s 272ms/step - loss: 0.0543 -
val_loss: 0.0593
Epoch 8/50
```

```

10/10 [=====] - 3s 378ms/step - loss: 0.0576 -
val_loss: 0.0607
Epoch 9/50
10/10 [=====] - 2s 269ms/step - loss: 0.0540 -
val_loss: 0.0533
Epoch 10/50
10/10 [=====] - 2s 272ms/step - loss: 0.0562 -
val_loss: 0.0663
Epoch 11/50
10/10 [=====] - 2s 261ms/step - loss: 0.0548 -
val_loss: 0.0615
Epoch 12/50
10/10 [=====] - 2s 265ms/step - loss: 0.0553 -
val_loss: 0.0574
Epoch 13/50
10/10 [=====] - 3s 372ms/step - loss: 0.0599 -
val_loss: 0.0703
Epoch 14/50
10/10 [=====] - 2s 271ms/step - loss: 0.0572 -
val_loss: 0.0601
Epoch 15/50
10/10 [=====] - 2s 267ms/step - loss: 0.0539 -
val_loss: 0.0583
Epoch 16/50
10/10 [=====] - 2s 269ms/step - loss: 0.0549 -
val_loss: 0.0555
Epoch 17/50
10/10 [=====] - 3s 300ms/step - loss: 0.0554 -
val_loss: 0.0605
Epoch 18/50
10/10 [=====] - 3s 325ms/step - loss: 0.0557 -
val_loss: 0.0652
Epoch 19/50
10/10 [=====] - 2s 273ms/step - loss: 0.0576 -
val_loss: 0.0615
Epoch 20/50
10/10 [=====] - 2s 267ms/step - loss: 0.0592 -
val_loss: 0.0725
Epoch 21/50
10/10 [=====] - 2s 266ms/step - loss: 0.0537 -
val_loss: 0.0594
Epoch 22/50
10/10 [=====] - 3s 359ms/step - loss: 0.0564 -
val_loss: 0.0631
Epoch 23/50
10/10 [=====] - 3s 273ms/step - loss: 0.0539 -
val_loss: 0.0614
Epoch 24/50

```

10/10 [=====] - 2s 271ms/step - loss: 0.0546 -  
val\_loss: 0.0701  
Epoch 25/50  
10/10 [=====] - 2s 271ms/step - loss: 0.0554 -  
val\_loss: 0.0694  
Epoch 26/50  
10/10 [=====] - 2s 269ms/step - loss: 0.0562 -  
val\_loss: 0.0714  
Epoch 27/50  
10/10 [=====] - 3s 383ms/step - loss: 0.0541 -  
val\_loss: 0.0487  
Epoch 28/50  
10/10 [=====] - 2s 270ms/step - loss: 0.0554 -  
val\_loss: 0.0553  
Epoch 29/50  
10/10 [=====] - 3s 273ms/step - loss: 0.0528 -  
val\_loss: 0.0626  
Epoch 30/50  
10/10 [=====] - 3s 277ms/step - loss: 0.0566 -  
val\_loss: 0.0610  
Epoch 31/50  
10/10 [=====] - 3s 309ms/step - loss: 0.0549 -  
val\_loss: 0.0627  
Epoch 32/50  
10/10 [=====] - 3s 313ms/step - loss: 0.0522 -  
val\_loss: 0.0500  
Epoch 33/50  
10/10 [=====] - 2s 267ms/step - loss: 0.0547 -  
val\_loss: 0.0597  
Epoch 34/50  
10/10 [=====] - 3s 301ms/step - loss: 0.0530 -  
val\_loss: 0.0632  
Epoch 35/50  
10/10 [=====] - 3s 276ms/step - loss: 0.0547 -  
val\_loss: 0.0556  
Epoch 36/50  
10/10 [=====] - 3s 386ms/step - loss: 0.0533 -  
val\_loss: 0.0614  
Epoch 37/50  
10/10 [=====] - 2s 270ms/step - loss: 0.0558 -  
val\_loss: 0.0610  
Epoch 38/50  
10/10 [=====] - 3s 276ms/step - loss: 0.0538 -  
val\_loss: 0.0601  
Epoch 39/50  
10/10 [=====] - 2s 272ms/step - loss: 0.0509 -  
val\_loss: 0.0677  
Epoch 40/50

```

10/10 [=====] - 3s 281ms/step - loss: 0.0547 -
val_loss: 0.0638
Epoch 41/50
10/10 [=====] - 3s 346ms/step - loss: 0.0568 -
val_loss: 0.0588
Epoch 42/50
10/10 [=====] - 2s 262ms/step - loss: 0.0553 -
val_loss: 0.0622
Epoch 43/50
10/10 [=====] - 2s 265ms/step - loss: 0.0537 -
val_loss: 0.0583
Epoch 44/50
10/10 [=====] - 3s 298ms/step - loss: 0.0545 -
val_loss: 0.0680
Epoch 45/50
10/10 [=====] - 3s 351ms/step - loss: 0.0532 -
val_loss: 0.0650
Epoch 46/50
10/10 [=====] - 3s 273ms/step - loss: 0.0504 -
val_loss: 0.0594
Epoch 47/50
10/10 [=====] - 2s 271ms/step - loss: 0.0531 -
val_loss: 0.0611
Epoch 48/50
10/10 [=====] - 2s 272ms/step - loss: 0.0517 -
val_loss: 0.0567
Epoch 49/50
10/10 [=====] - 2s 272ms/step - loss: 0.0553 -
val_loss: 0.0648
Epoch 50/50
10/10 [=====] - 3s 381ms/step - loss: 0.0525 -
val_loss: 0.0546

```

### Visualize learning process

```

[ ]: plt.plot(history.history['loss'], label = 'Train', c = 'b')
plt.plot(history.history['val_loss'], label = 'Validation', c = 'r')
plt.legend()
plt.xlabel('Epoch', fontsize = 16)
plt.ylabel('Loss', fontsize = 16)


```

```

[ ]: Text(0, 0.5, 'Loss')

```





```
[ ]: for pair, y in generator(X_test, y_test):
    y_pred = model.predict(pair)
    print('Batch_size: ', len(y))
    idx = np.random.choice(range(len(y)))
    print('Pair label:', y[idx])
    print('Distance:', y_pred[idx])

    f1 = cnn(pair[0])
    f2 = cnn(pair[1])
    d = np.sqrt(np.sum((f1-f2)**2,axis = 1, keepdims = True))
    print('Distance by features:', d[idx])


    plt.subplot(121)
    plt.imshow(pair[0][idx].reshape(28,28),cmap = 'gray')
    plt.subplot(122)
    plt.imshow(pair[1][idx].reshape(28,28),cmap = 'gray')
    break
```

5/5 [=====] - 0s 3ms/step

Batch\_size: 160

Pair label: 1.0

Distance: [0.31551546]  
Distance by features: [0.31551543]




**Visualize new feature space**

```
[ ]: f = cnn.predict(X_test/255.)  
p = plt.scatter(f[:,0],f[:,1], c = y_test, s=1)  
plt.colorbar(p)
```

313/313 [=====] - 1s 2ms/step

```
[ ]: <matplotlib.colorbar.Colorbar at 0x7f9a280bdcf0>
```



### Save model

```
[ ]: cnn.save('cnn_loss1.h5')
```

WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile\_metrics` will be empty until you train or evaluate the model.

### Load model and test


```
[ ]: from tensorflow.keras.models import load_model
m = load_model('cnn_loss1.h5')

f1 = m.predict(X_test / 255.)
p = plt.scatter(f1[:,0],f1[:,1],c=y_test,s=1)
plt.colorbar(p)
```

WARNING:tensorflow:No training configuration found in the save file, so the model was \*not\* compiled. Compile it manually.

313/313 [=====] - 1s 1ms/step

```
[ ]: <matplotlib.colorbar.Colorbar at 0x7f9a207dbd60>
```



Visualize negative distance and positive distance

```
[ ]: i = 0
y_true = []
y_pred = []
for pair,y in generator(X_test,y_test):
    f1 = cnn(pair[0])
    f2 = cnn(pair[1])
    d = np.sqrt(np.sum((f1 - f2)**2, axis = 1, keepdims = True))
    y_pred +=list(d.ravel())
    y_true +=list(y)
    i+=1
    if i>500:
        break
```


Plot histogram

```
[ ]: y_pred = np.array(y_pred)
y_true = np.array(y_true)

positive_distances = y_pred[y_true == 1]
negative_distances = y_pred[y_true == 0]
```

```
plt.hist(positive_distances, color = 'r', density = True, bins = 20)
plt.hist(negative_distances, color = 'b', density = True, bins = 20)
```

```
[ ]: (array([1.52630560e-02, 5.08190388e-02, 1.32511077e-01, 3.50703401e-01,
        6.85276668e-01, 8.65484474e-01, 8.56118864e-01, 6.00636067e-01,
        5.61263847e-01, 7.08171257e-01, 6.38967159e-01, 4.40894277e-01,
        3.31798093e-01, 2.73347514e-01, 2.08652784e-01, 1.34765763e-01,
        6.98979312e-02, 2.23742387e-02, 4.16265595e-03, 5.20331132e-04]),
array([0.02938122, 0.17323241, 0.3170836 , 0.46093479, 0.60478598,
        0.74863714, 0.89248836, 1.03633952, 1.18019068, 1.32404196,
        1.46789312, 1.61174428, 1.75559545, 1.89944661, 2.04329777,
        2.18714905, 2.33100033, 2.47485137, 2.61870265, 2.76255369,
        2.90640497])),
<BarContainer object of 20 artists>)
```



Check report using sklearn

```
[ ]: thresh = 0.5
y_pred_ = y_pred < thresh
y_pred_.astype('uint8')
from sklearn.metrics import classification_report
```

```
print(classification_report(y_true, y_pred_))
```

	precision	recall	f1-score	support
0.0	0.93	0.96	0.94	40080
1.0	0.96	0.92	0.94	40080
accuracy			0.94	80160
macro avg	0.94	0.94	0.94	80160
weighted avg	0.94	0.94	0.94	80160